Completed generation comparisons: fixed 40 cases

Seed 42; native text guidance 5; active physical guidance 5.

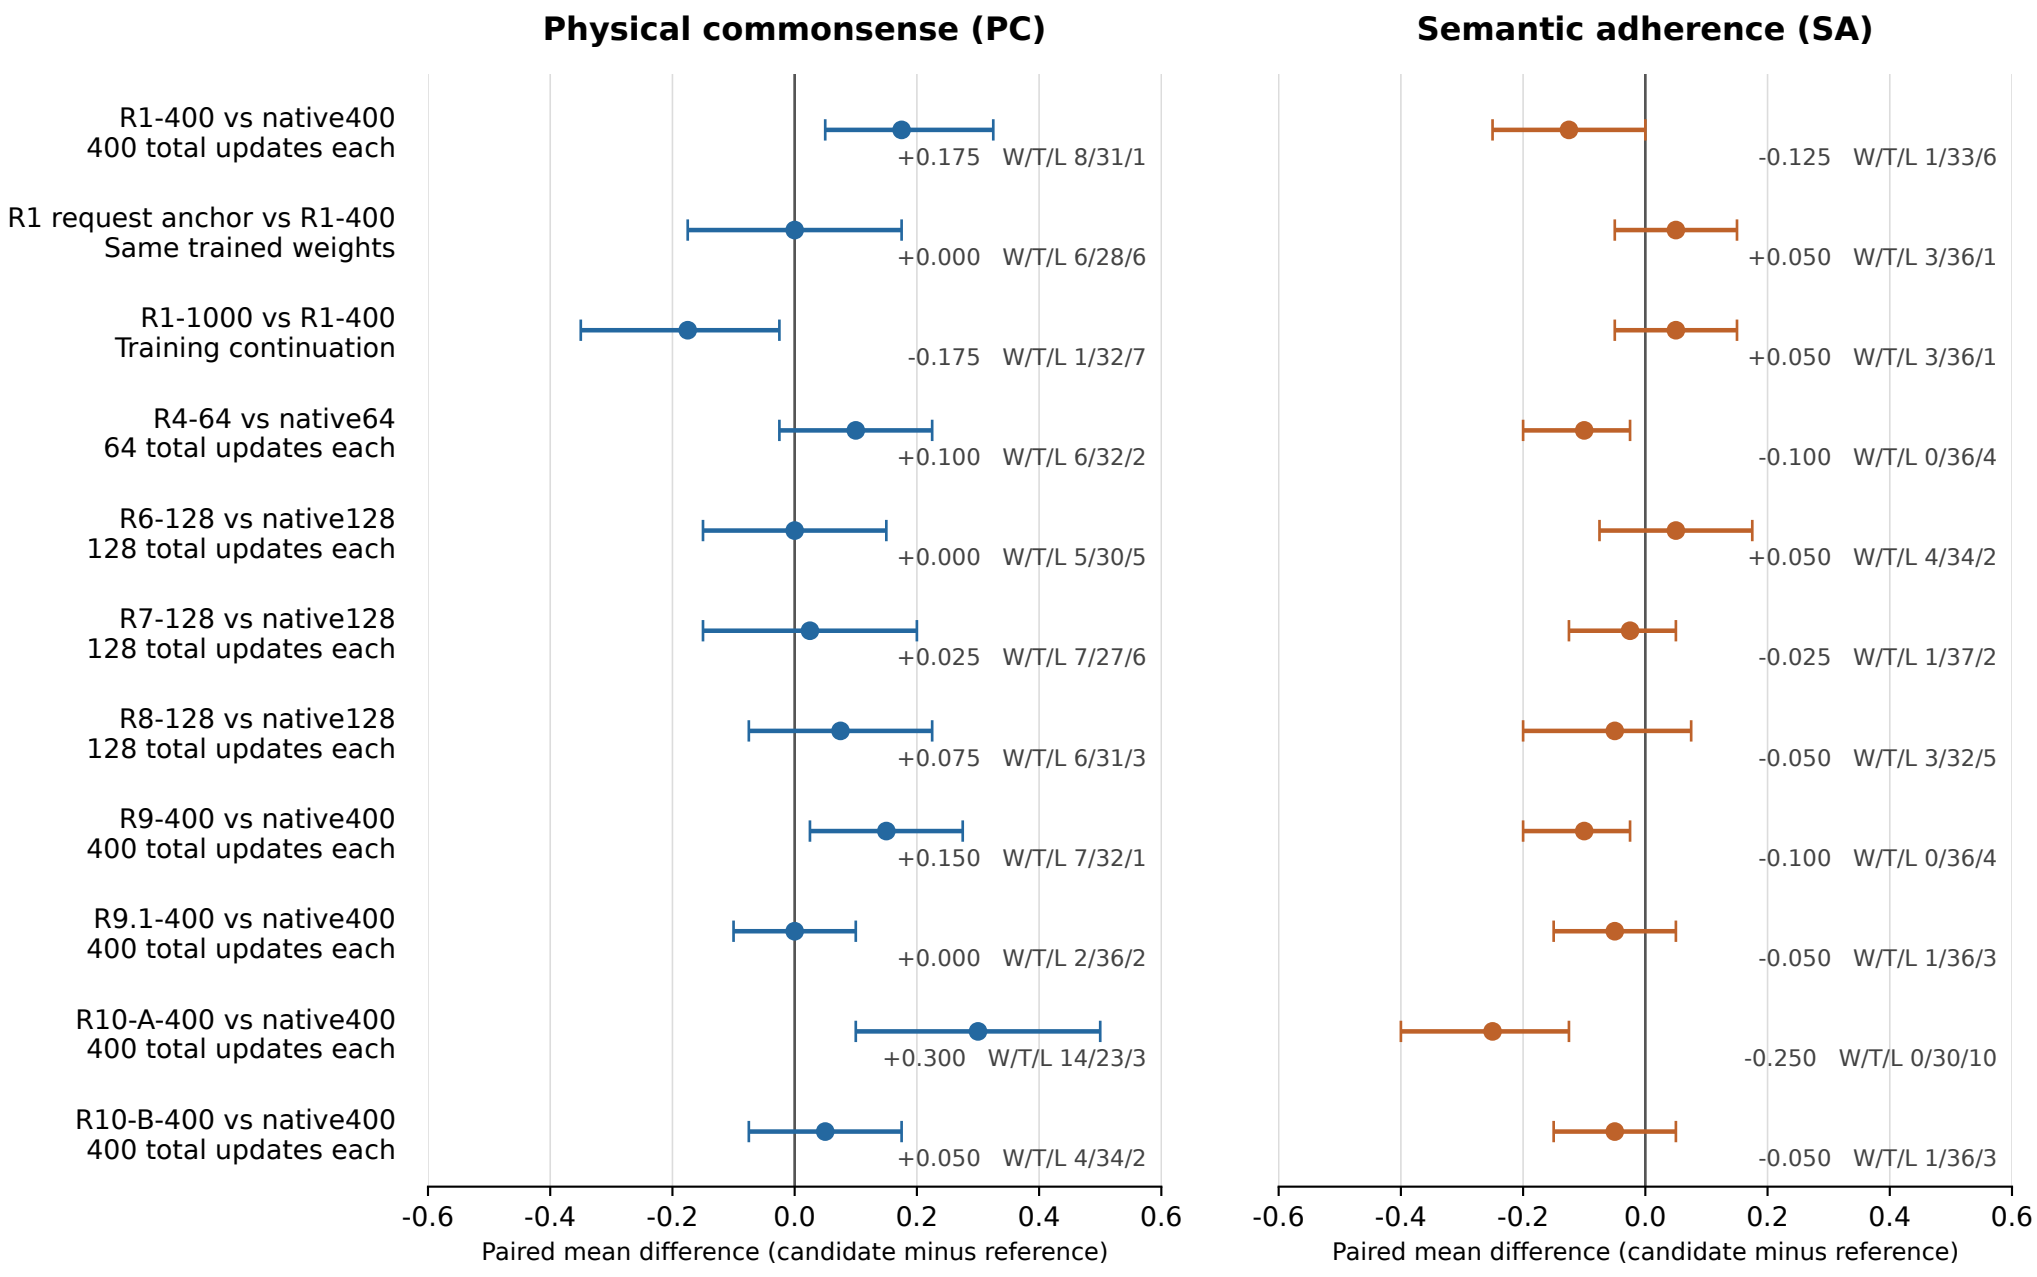

Units: points on the 1–5 scale. Lines: existing case-paired bootstrap 95% intervals; they do not estimate variation across seeds.
R1/R9/R9.1-400:128 interface+272 LoRA; R4-64:16+48; R6/R7/R8-128:16+112. Native baselines train LoRA on every update.
The continuation row deliberately compares different training budgets. These are shared-first-image I2V results, not the full T2V benchmark.